Exercise 5

Machine Learning I

4A-1.

We know:

$$p(\theta) = \frac{1}{\Delta prior}$$

$$p(\theta_{MAP}|D) \approx \frac{1}{\Delta posterior}$$

Using Bayes' theorem and rearranging terms we get:

$$\begin{aligned} p(\theta|D) &= \frac{p(D|\theta)p(\theta)}{p(D)} \\ (*) \quad \approx \quad p(D) &= p(D|\theta_{MAP})\frac{\Delta posterior}{\Delta prior} \end{aligned}$$

The approximately "≈" is valid because most density is concentrated around θ_{Map}.
As we saw in the lecture slides, the evidence approximation is useful to make inferences about model complexity and fit.

Let us see what happens in $(*)$ if we vary the input D:

As the likelihood $p(D|\theta)$ is strongly related to the fit, $p(D)$ acquires information about the model fit. As fit is related to model complexity, a better fit usually entails a more complex model as well. Additionally, because of $\frac{\Delta posterior}{\Delta prior}$ our $p(D)$ also receives information about model complexity and plausibility.

If $\Delta posterior$ is too high compared to our belief $\Delta prior$, then $p(D)$ goes down. Think about what this means. A very high posterior around θ_{MAP} leads to a point estimate not unlike the maximum likelihood method. This may lead to little bias in the input set $\mathbf{x}_i$. As we know from the bias/variance decomposition, this may increase variance for different datasets $\mathbf{x}_j$. Which means we could have learned too much from our data.

Because $p(D)$ goes down in this case, it helps us combat overfitting.

On the other hand, if our fit is good ($p(D|\theta_{MAP})$ high), yet we have a general enough model ($p(\theta_{MAP}|D)$ low), we receive a good value for $p(D)$.

The optimal value for $p(D)$ is thereby a balance of model complexity, belief ($\frac{\Delta posterior}{\Delta prior}$) and fit ($p(D|\theta_{MAP})$).

See more on page 163, Bishop PRML.

4A-2.

Prerequisites

Let $\mathbf{A}$ be invertible, $\mathbf{B}, \mathbf{C}$ arbitrary and $\mathbf{I}$ be the identity matrix.

We then have:

$$(*) \quad (\mathbf{A}+\mathbf{BCD})^{-1} = \mathbf{A}^{-1} - (\mathbf{I}+\mathbf{A}^{-1}\mathbf{BCD})^{-1}\mathbf{A}^{-1}\mathbf{BCDA}^{-1}$$

The proof of above can be seen here:

http://www0.cs.ucl.ac.uk/staff/gridgway/mil/mil.pdf

Also, as seen in the last lecture slides, the posterior of Bayesian linear regression with normal prior is:

$$p(\mathbf{w}|\mathbf{t}, \mathbf{\Sigma}_N, \tau_e) = N(\mathbf{w}; \boldsymbol{\mu}_N, \mathbf{\Sigma}_N)$$

$$\boldsymbol{\mu}_N = \tau_e \mathbf{\Sigma}_N \mathbf{\Phi}^T \mathbf{t}$$

$$\mathbf{\Sigma}_N = \left(\Sigma_0^{-1} + \tau_e \mathbf{\Phi}^{\mathrm{T}} \mathbf{\Phi}\right)^{-1}$$

Let us have K old and M "new" datapoints.
Because the predictive distribution is dependent on the posterior, it makes sense to calculate the new posterior first.

As seen in the task, we have:

$$p(\mathbf{w}|\mathbf{t}_M, \mathbf{t}_K) \propto p(\mathbf{w}|\mathbf{t}_K) p(\mathbf{t}_M|\mathbf{w})$$

This leads to:

$$p(\mathbf{w}|\mathbf{t}_K) p(\mathbf{t}_M|\mathbf{w}) \quad \propto e^{-\frac{1}{2}[(\mathbf{w}-\boldsymbol{\mu}_K)^T \mathbf{\Sigma}_K^{-1} (\mathbf{w}-\boldsymbol{\mu}_K) + (\mathbf{t}_m - \mathbf{\Phi}_M \mathbf{w})^T \tau_e \mathbf{I} (\mathbf{t}_m - \mathbf{\Phi}_M \mathbf{w})]}$$

Note: We have $\tau_e \mathbf{I}$ because the predictions are independent from each other. Focusing on the exponent:

$$
\begin{aligned}
(\mathbf{w}-\boldsymbol{\mu}_K)^T\boldsymbol{\Sigma}_K^{-1}(\mathbf{w}-\boldsymbol{\mu}_K)+(\mathbf{t}_M-\boldsymbol{\Phi}_M\mathbf{w})^T\tau_e\mathbf{I}(\mathbf{t}_M-\boldsymbol{\Phi}_M\mathbf{w}) &= \mathbf{w}^T\underbrace{(\boldsymbol{\Sigma}_K^{-1}+\tau_e\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)}_{\boldsymbol{\Sigma}_{K+M}^{-1}}\mathbf{w}-2\boldsymbol{w}^T[\boldsymbol{\Sigma}_K^{-1}\boldsymbol{\mu}_K+\tau_e\boldsymbol{\Phi}_M^T\mathbf{t}_M] \\
&\quad +\boldsymbol{\mu}_K\boldsymbol{\Sigma}_K^{-1}\boldsymbol{\mu}_M+\tau_e\boldsymbol{t}_M^T\boldsymbol{t}_M \\
&= \mathbf{w}^T\boldsymbol{\Sigma}_{K+M}^{-1}\mathbf{w} \\
&\quad -2\boldsymbol{w}^T\underbrace{\boldsymbol{\Sigma}_{K+M}^{-1}\boldsymbol{\Sigma}_{K+M}}_{=1}[\boldsymbol{\Sigma}_K^{-1}\boldsymbol{\mu}_K+\tau_e\boldsymbol{\Phi}_M^T\mathbf{t}_M] \\
&\quad +c
\end{aligned}
$$

This allows us to complete the square above and we get the updated posterior:

$$p(\mathbf{w}|\mathbf{t}_M,\mathbf{t}_K)=N(\mathbf{w};\boldsymbol{\mu}_{M+K},\boldsymbol{\Sigma}_{K+M})$$

with

$$\boldsymbol{\Sigma}_{K+M}=(\boldsymbol{\Sigma}_K^{-1}+\tau_e\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}$$

$$\boldsymbol{\mu}_{M+K}=\boldsymbol{\Sigma}_{K+M}[\boldsymbol{\Sigma}_K^{-1}\boldsymbol{\mu}_K+\tau_e\boldsymbol{\Phi}_M^T\mathbf{t}_M]$$

Note: When $\boldsymbol{\mu}_K=\mathbf{0}$ we get the same result as in the slides for the prior $p(\mathbf{w}|\boldsymbol{\Sigma}_0)$.

Now we use $(*)$ for the covariance:

$$(\boldsymbol{\Sigma}_K^{-1}+\tau_e\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}=\boldsymbol{\Sigma}_K-\tau_e(\mathbf{I}+\tau_e\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M\boldsymbol{\Sigma}_K$$

As seen in the slides (or the task), we try to show:

$$\sigma_{K+M}(x^*)\leq\sigma_K(x^*)$$

where

$$\sigma_n=\frac{1}{r_e}+\boldsymbol{\phi}(x^*)^T\boldsymbol{\Sigma}_N\boldsymbol{\phi}(x^*)$$

Now we just plug our new covariance matrices into σ_n:

$$
\begin{aligned}
\sigma_{K+M}(x^*) &= \frac{1}{r_e}\boldsymbol{\phi}(x^*)^T(\boldsymbol{\Sigma}_K^{-1}+\tau_e\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}\boldsymbol{\phi}(x^*) \\
&= \frac{1}{r_e}\boldsymbol{\phi}(x^*)^T[\boldsymbol{\Sigma}_K-\tau_e(\mathbf{I}+\tau_e\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M\boldsymbol{\Sigma}_K]\boldsymbol{\phi}(x^*) \\
&= \underbrace{\frac{1}{r_e}\boldsymbol{\phi}(x^*)^T\boldsymbol{\Sigma}_K\boldsymbol{\phi}(x^*)}_{\sigma_K(x^*)}-\frac{1}{r_e}\boldsymbol{\phi}(x^*)^T[\tau_e(\mathbf{I}+\tau_e\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M\boldsymbol{\Sigma}_K]\boldsymbol{\phi}(x^*) \\
&\leq \sigma_K(x^*).
\end{aligned}
$$

The last equality holds because

$$\tau_e(\mathbf{I}+\tau_e\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M)^{-1}\boldsymbol{\Sigma}_K\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M\boldsymbol{\Sigma}_K$$

is at least positive semidefinite (for why, look up the properties of positive semidefinite matrices, especially in regard to $\boldsymbol{\Sigma}_K$ and $\boldsymbol{\Phi}_M^T\boldsymbol{\Phi}_M$.

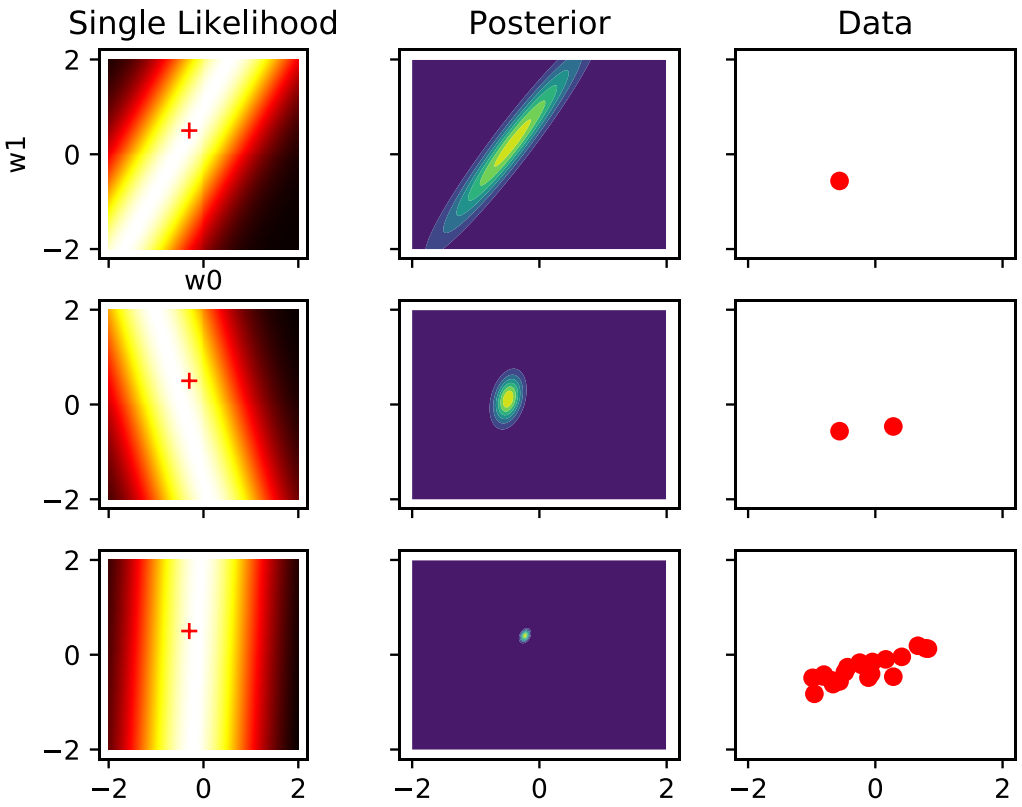

Figure 1 Replication of Bishop's Figure 3.7 in Python. The Likelihood is of a single data point

4A-3.

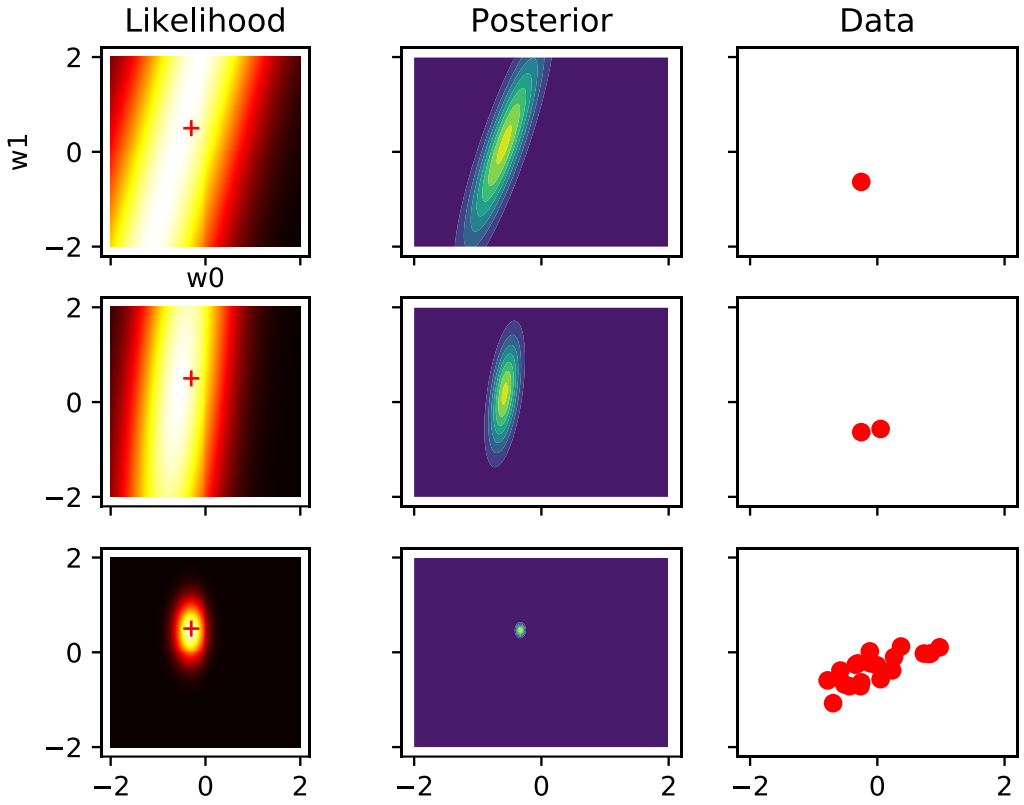

Figure 2 Replication of Bishop's Figure 3.7 in Python this time with joint likelihood of the samples.